

# Data-Classifiers

December 17, 2022

## 1 Problem Statement

Given the MAGIC gamma telescope dataset that can be obtained using this [Link](#). This dataset is generated to simulate registration of high energy gamma particles in a ground-based atmospheric Cherenkov gamma telescope using the imaging technique. The dataset consists of two classes; gammas (signal) and hadrons (background). There are 12332 gamma events and 6688 hadron events. You are required to use this dataset to apply different classification models such as Decision Trees, Naive Bayes Classifier, Random Forests, AdaBoost and K-Nearest Neighbor (K-NN). You are also required to tune the parameters of these models, and compare the performance of models with each other.

### 1.1 Attribute Information:

1. fLength: continuous # major axis of ellipse [mm]
2. fWidth: continuous # minor axis of ellipse [mm]
3. fSize: continuous # 10-log of sum of content of all pixels [in #phot]
4. fConc: continuous # ratio of sum of two highest pixels over fSize [ratio]
5. fConc1: continuous # ratio of highest pixel over fSize [ratio]
6. fAsym: continuous # distance from highest pixel to center, projected onto major axis [mm]
7. fM3Long: continuous # 3rd root of third moment along major axis [mm]
8. fM3Trans: continuous # 3rd root of third moment along minor axis [mm]
9. fAlpha: continuous # angle of major axis with vector to origin [deg]
10. fDist: continuous # distance from origin to center of ellipse [mm]
11. class: g,h # gamma (signal), hadron (background)

g = gamma (signal): 12332 h = hadron (background): 6688

### 1.2 Importing libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
```

### 1.3 Importing the Dataset

```
[2]: # Adding column names
col_names = ['fLength', 'fWidth', 'fSize', 'fConc', 'fConc1', 'fAsym', 'fM3Long', 'fM3Trans', 'fAlpha', 'fDist', 'class']
# Importing the dataset
df = pd.read_csv('magic04.data', names = col_names)
#df = df.sample(frac = 1)
df.head()
```

```
[2]:
```

	fLength	fWidth	fSize	fConc	fConc1	fAsym	fM3Long	fM3Trans	\
0	28.7967	16.0021	2.6449	0.3918	0.1982	27.7004	22.0110	-8.2027	
1	31.6036	11.7235	2.5185	0.5303	0.3773	26.2722	23.8238	-9.9574	
2	162.0520	136.0310	4.0612	0.0374	0.0187	116.7410	-64.8580	-45.2160	
3	23.8172	9.5728	2.3385	0.6147	0.3922	27.2107	-6.4633	-7.1513	
4	75.1362	30.9205	3.1611	0.3168	0.1832	-5.5277	28.5525	21.8393	

	fAlpha	fDist	class
0	40.0920	81.8828	g
1	6.3609	205.2610	g
2	76.9600	256.7880	g
3	10.4490	116.7370	g
4	4.6480	356.4620	g

```
[3]: print(df['class'].value_counts())
df.describe()
```

```
g    12332
h     6688
Name: class, dtype: int64
```

```
[3]:
```

	fLength	fWidth	fSize	fConc	fConc1	\
count	19020.000000	19020.000000	19020.000000	19020.000000	19020.000000	
mean	53.250154	22.180966	2.825017	0.380327	0.214657	
std	42.364855	18.346056	0.472599	0.182813	0.110511	
min	4.283500	0.000000	1.941300	0.013100	0.000300	

25%	24.336000	11.863800	2.477100	0.235800	0.128475
50%	37.147700	17.139900	2.739600	0.354150	0.196500
75%	70.122175	24.739475	3.101600	0.503700	0.285225
max	334.177000	256.382000	5.323300	0.893000	0.675200

	fAsym	fM3Long	fM3Trans	fAlpha	fDist
count	19020.000000	19020.000000	19020.000000	19020.000000	19020.000000
mean	-4.331745	10.545545	0.249726	27.645707	193.818026
std	59.206062	51.000118	20.827439	26.103621	74.731787
min	-457.916100	-331.780000	-205.894700	0.000000	1.282600
25%	-20.586550	-12.842775	-10.849375	5.547925	142.492250
50%	4.013050	15.314100	0.666200	17.679500	191.851450
75%	24.063700	35.837800	10.946425	45.883550	240.563825
max	575.240700	238.321000	179.851000	90.000000	495.561000

## 1.4 Data Balancing

```
[4]: # Splitting dataset by class label
df_g = df[df['class'] == 'g']
df_h = df[df['class'] == 'h']
df_g
```

```
[4]:
```

	fLength	fWidth	fSize	fConc	fConc1	fAsym	fM3Long	\
0	28.7967	16.0021	2.6449	0.3918	0.1982	27.7004	22.0110	
1	31.6036	11.7235	2.5185	0.5303	0.3773	26.2722	23.8238	
2	162.0520	136.0310	4.0612	0.0374	0.0187	116.7410	-64.8580	
3	23.8172	9.5728	2.3385	0.6147	0.3922	27.2107	-6.4633	
4	75.1362	30.9205	3.1611	0.3168	0.1832	-5.5277	28.5525	
...	...	...	...	...	...	...	...	
12327	12.8703	11.4444	2.3811	0.7360	0.3805	-15.0946	5.3032	
12328	26.8595	20.5946	2.8754	0.3438	0.2152	-3.4556	-20.0014	
12329	22.0913	10.8949	2.2945	0.5381	0.2919	15.2776	18.2296	
12330	56.2216	18.7019	2.9297	0.2516	0.1393	96.5758	-41.2969	
12331	31.5125	19.2867	2.9578	0.2975	0.1515	38.1833	21.6729	

	fM3Trans	fAlpha	fDist	class
0	-8.2027	40.0920	81.8828	g
1	-9.9574	6.3609	205.2610	g
2	-45.2160	76.9600	256.7880	g
3	-7.1513	10.4490	116.7370	g
4	21.8393	4.6480	356.4620	g
...	...	...	...	...
12327	11.6208	21.0120	204.0370	g
12328	-9.0535	3.9848	205.4980	g
12329	7.3975	21.0680	123.2810	g
12330	11.3764	5.9110	197.2090	g
12331	-12.0726	17.5809	171.2270	g

[12332 rows x 11 columns]

```
[5]: # Balancing Data
#dfbalanced_g = df_g.sample(df_h.shape[0])
dfbalanced_g = df_g[:df_h.shape[0]]
df = pd.concat([dfbalanced_g, df_h], axis=0) # concatenate horizontally
df['class'].value_counts()
```

```
[5]: g    6688
     h    6688
     Name: class, dtype: int64
```

```
[6]: # Separating features and class
X = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
y
```

```
[6]: array(['g', 'g', 'g', ..., 'h', 'h', 'h'], dtype=object)
```

## 1.5 Encoding class labels

```
[7]: # Encoding the Dependent Variable(y)
from sklearn.preprocessing import LabelEncoder
le= LabelEncoder()
y = le.fit_transform(y)
print(y)
```

```
[0 0 0 ... 1 1 1]
```

## 1.6 Splitting Dataset into Train set and Test set

```
[12]: from sklearn.model_selection import train_test_split
      from collections import Counter
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2999,
      ↪random_state = 0, stratify=y)
      print(f"Train set size: {X_train.shape[0]}\nTest set size: {X_test.shape[0]}")
      print(Counter(y_train))
      print(Counter(y_test))
```

Train set size: 9364

Test set size: 4012

Counter({1: 4682, 0: 4682})

Counter({0: 2006, 1: 2006})

## 1.7 Feature Scaling

```
[13]: # for KNN as it is based upon distances
from sklearn.preprocessing import MinMaxScaler
mm = MinMaxScaler()
X_train = mm.fit_transform(X_train)
X_test = mm.transform(X_test)
print(X_train)
print(X_test)

[[0.03866115 0.05030969 0.11402557 ... 0.51384336 0.61213889 0.33926103]
 [0.1060195  0.04611322 0.19295707 ... 0.46895876 0.85389222 0.47096393]
 [0.18285673 0.10055191 0.41621442 ... 0.52365604 0.03552222 0.34721202]
 ...
 [0.0334423  0.03998955 0.11643753 ... 0.46984637 0.58271778 0.44295361]
 [0.0832623  0.05837851 0.13835625 ... 0.45035779 0.89218333 0.14212638]
 [0.04340349 0.02731705 0.01112518 ... 0.47108938 0.25136667 0.28382466]]
[[0.18817527 0.08000094 0.27695369 ... 0.45629629 0.10407333 0.30876203]
 [0.1695148  0.08601579 0.41292812 ... 0.519218  0.02028444 0.42572647]
 [0.25432159 0.11695868 0.41983237 ... 0.55411699 0.04223333 0.52288427]
 ...
 [0.05586736 0.06523547 0.14363242 ... 0.44734942 0.52550222 0.31418812]
 [0.067696   0.06311442 0.16374216 ... 0.51332161 0.43543333 0.33746447]
 [0.05231504 0.0401678  0.13262783 ... 0.47074563 0.02010222 0.49621104]]
```

## 1.8 Model Creation and Evaluation Functions

```
[14]: # Model Creation
#from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import f1_score, \
    accuracy_score, recall_score, precision_score
def train_model(X_train, X_test, y_train, classifier):
    classifier.fit(X_train, y_train)
    pred = classifier.predict(X_test)
    return pred

[15]: comparison_cols=["Classifier", "Accuracy", "F1 score", "Precision", "Recall"]
comparison_df = pd.DataFrame(columns=comparison_cols)

[16]: def evaluate(name, y_pred, y_test):
    global comparison_df
    f1 = f1_score(y_test, y_pred)
    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    print(f"{name} accuracy: {acc}")
    print(f"{name} f1 score: {f1}")
```

```

print(f"{name} precision: {prec}")
print(f"{name} recall: {recall}")
entry = pd.DataFrame([[name, acc*100, f1*100, prec*100, recall*100]],
↳columns=comparison_cols)
comparison_df = comparison_df.append(entry)
#pd.concat([comparison_df, entry])

```

## 1.9 Hyperparameter tuning function

```

[17]: from sklearn.model_selection import StratifiedKFold
def cross_validation(X_train, y_train, classifier, cv=10):
    # stratified k-fold is used as the folds are made by preserving the
↳percentage of samples for each class.
    skf = StratifiedKFold(n_splits=cv)
    avg_accuracy = 0
    sum_accuracy = 0
    for train_index, test_index in skf.split(X_train, y_train):
        x_train_folds = X_train[train_index]
        y_train_folds = y_train[train_index]
        x_test_fold = X_train[test_index]
        y_test_fold = y_train[test_index]
        classifier.fit(x_train_folds, y_train_folds)
        y_pred_fold = classifier.predict(x_test_fold)
        sum_accuracy += accuracy_score(y_pred_fold, y_test_fold)
    avg_accuracy = sum_accuracy/cv
    return avg_accuracy

```

## 1.10 Decision Tree Classifier

Tunable parameters: None

```

[18]: y_pred =
↳train_model(X_train, X_test, y_train, DecisionTreeClassifier(random_state=5))
evaluate("Decision Tree", y_pred, y_test)

```

Decision Tree accuracy: 0.790877367896311  
Decision Tree f1 score: 0.7878634639696587  
Decision Tree precision: 0.7993842996408415  
Decision Tree recall: 0.7766699900299102

/tmp/ipykernel\_103467/1172794529.py:12: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
comparison_df = comparison_df.append(entry)
```

```

[19]: from sklearn.metrics import confusion_matrix
cf_matrix = confusion_matrix(y_test, y_pred)

```

```


print(cf_matrix)
ax= plt.subplot()
sns.heatmap(cf_matrix, annot=True,fmt='g',ax=ax,cmap='Blues')
ax.set_xlabel('Predicted labels');ax.set_ylabel('Actual labels');
ax.set_title('Confusion Matrix');
ax.xaxis.set_ticklabels(['Gamma', 'Hadron']); ax.yaxis.set_ticklabels(['Gamma',
↪ 'Hadron']);

```

```

[[1615  391]
 [ 448 1558]]

```



### 1.11 AdaBoost Classifier

An AdaBoost classifier is a meta-estimator that begins by fitting a classifier on the original dataset and then fits additional copies of the classifier on the same dataset but where the weights of incorrectly classified instances are adjusted such that subsequent classifiers focus more on difficult cases.

tunable parameters: **n\_estimators**

**n\_estimators**: The maximum number of estimators at which boosting is terminated. In case of

perfect fit, the learning procedure is stopped early.


## 1.12 Hyperparameter tuning

```
[20]: import time
# possible values of n_estimators
n_trees = np.linspace(100,500,10,dtype=int)
avg_accuracies = []
for n in n_trees:
    adaboost = AdaBoostClassifier(n_estimators = n)
    t1 = time.time()
    acc_score = cross_validation(X_train,y_train,adaboost,cv=10)
    t2 = time.time() - t1
    print(f"n = {n}\naverage accuracy = {acc_score}\ntime = {t2}")
    avg_accuracies.append(acc_score)
```

n = 100	average accuracy = 0.8221931468863165	time = 10.614821434020996
n = 144	average accuracy = 0.8212330906968047	time = 15.230633735656738
n = 188	average accuracy = 0.8237942287168541	time = 19.895525217056274
n = 233	average accuracy = 0.8224063660162912	time = 24.690342903137207
n = 277	average accuracy = 0.8218729761285791	time = 29.395816326141357
n = 322	average accuracy = 0.8217659104798912	time = 34.204378604888916
n = 366	average accuracy = 0.8221928048235412	time = 39.30795097351074
n = 411	average accuracy = 0.8229395278621532	time = 46.42223405838013
n = 455	average accuracy = 0.823367334373204	time = 52.46410298347473
n = 500	average accuracy = 0.8216582747265779	time = 54.33817172050476

```
[22]: plt.plot(n_trees,np.array(avg_accuracies))
plt.show()
```





```
[23]: n_trees = np.linspace(150,250,20,dtype=int)
avg_accuracies = []
for n in n_trees:
    adaboost = AdaBoostClassifier(n_estimators = n)
    t1 = time.time()
    acc_score = cross_validation(X_train,y_train,adaboost,cv=10)
    t2 = time.time() - t1
    print(f"n = {n}\naverage accuracy = {acc_score}\ntime = {t2}")
    avg_accuracies.append(acc_score)
```

```
n = 150 average accuracy = 0.8220874494887301    time = 16.025478839874268
n = 155 average accuracy = 0.8229414662178802    time = 16.452531814575195
n = 160 average accuracy = 0.8252899552125805    time = 17.001068592071533
n = 165 average accuracy = 0.8250765080407556    time = 17.56627869606018
n = 171 average accuracy = 0.8257176477027064    time = 18.1414852142334
n = 176 average accuracy = 0.8231535451386037    time = 18.66496992111206
n = 181 average accuracy = 0.823580667524104     time = 19.3525652885437
n = 186 average accuracy = 0.824007675888679     time = 19.782296180725098
n = 192 average accuracy = 0.8241141714327413    time = 20.394627332687378
n = 197 average accuracy = 0.8235804394822537    time = 21.102035999298096
n = 202 average accuracy = 0.8231536591595289    time = 21.463186979293823
```

```


n = 207 average accuracy = 0.8233673343732042    time = 21.96934413909912
n = 213 average accuracy = 0.8233673343732042    time = 22.6165874004364
n = 218 average accuracy = 0.8233669923104289    time = 23.104921579360962
n = 223 average accuracy = 0.8242214651232794    time = 23.70276951789856
n = 228 average accuracy = 0.8229402119877038    time = 24.56711745262146
n = 234 average accuracy = 0.8214452836384533    time = 25.188450813293457
n = 239 average accuracy = 0.8212321785294037    time = 25.73282241821289
n = 244 average accuracy = 0.8218720639611782    time = 30.268824815750122
n = 250 average accuracy = 0.8218724060239534    time = 29.331119537353516

```

```

[24]: plt.plot(n_trees,np.array(avg_accuracies))
      plt.show()

```



```

[35]: n_trees = np.append(np.arange(160,180,2), np.arange(210,230,2))
      avg_accuracies = []
      for n in n_trees:
          adaboost = AdaBoostClassifier(n_estimators = n)
          t1 = time.time()
          acc_score = cross_validation(X_train,y_train,adaboost,cv=10)
          t2 = time.time() - t1
          print(f"n = {n}\naverage accuracy = {acc_score}\ntime = {t2}")
          avg_accuracies.append(acc_score)


```

n = 160	average accuracy = 0.8252899552125805	time = 20.920289039611816
n = 162	average accuracy = 0.8257169635771557	time = 22.278916835784912
n = 164	average accuracy = 0.825503744447181	time = 20.267263412475586
n = 166	average accuracy = 0.8248630608689306	time = 24.5304536819458
n = 168	average accuracy = 0.8246498417389561	time = 25.378095626831055
n = 170	average accuracy = 0.8250767360826059	time = 27.45654058456421
n = 172	average accuracy = 0.8249700124966933	time = 32.40224051475525
n = 174	average accuracy = 0.8257169635771557	time = 32.43246912956238
n = 176	average accuracy = 0.8231535451386037	time = 27.431811571121216
n = 178	average accuracy = 0.8231538872013792	time = 29.31094717979431
n = 210	average accuracy = 0.8227264227531036	time = 33.8264045715332
n = 212	average accuracy = 0.8225129755812786	time = 34.81118130683899
n = 214	average accuracy = 0.8225128615603534	time = 35.082154273986816
n = 216	average accuracy = 0.8217652263543405	time = 37.366798400878906
n = 218	average accuracy = 0.8233669923104289	time = 37.05272030830383
n = 220	average accuracy = 0.8232604967663665	time = 39.5594539642334
n = 222	average accuracy = 0.8236878471937169	time = 37.14353060722351
n = 224	average accuracy = 0.8233676764359794	time = 35.75477313995361
n = 226	average accuracy = 0.8233675624150543	time = 36.94122648239136
n = 228	average accuracy = 0.8229402119877038	time = 37.269956827163696

```
[36]: plt.plot(n_trees,np.array(avg_accuracies))

def annot_max(x,y, ax=None):
    xmax = x[np.argmax(y)]
    ymax = y.max()
    text= "x={:.3f}, y={:.3f}".format(xmax, ymax)
    if not ax:
        ax=plt.gca()
    bbox_props = dict(boxstyle="square,pad=0.3", fc="w", ec="k", lw=0.72)
    arrowprops=dict(arrowstyle="->",connectionstyle="angle,angleA=0,angleB=60")
    kw = dict(xycoords='data',textcoords="axes fraction",
              arrowprops=arrowprops, bbox=bbox_props, ha="right", va="top")
    ax.annotate(text, xy=(xmax, ymax), xytext=(0.94,0.96), **kw)

annot_max(n_trees,np.array(avg_accuracies))
plt.show()
```



Best n\_estimators value: ~162

### 1.13 Hyperparameter tuning using Randomized Search

```
[23]: '''from sklearn.model_selection import RepeatedStratifiedKFold
from sklearn.model_selection import RandomizedSearchCV
grid = dict()
grid['n_estimators'] = [100, 200, 211, 300]
grid['learning_rate'] = [0.001, 0.01, 0.1, 1.0]
# define the evaluation procedure
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
# define the grid search procedure
grid_search = RandomizedSearchCV(estimator=AdaBoostClassifier(),
    ↪ param_distributions=grid, cv=cv, scoring='accuracy')
# execute the grid search
grid_result = grid_search.fit(X_train, y_train)
# summarize the best score and configuration
print("Best: %f using %s" % (grid_result.best_score_, grid_result.
    ↪ best_params_))'''
```

```
[23]: 'from sklearn.model_selection import RepeatedStratifiedKFold\nfrom
sklearn.model_selection import RandomizedSearchCV\ngrid =
```

```
dict()\ngrid[\n'n_estimators\'] = [100, 200, 211,300]\ngrid[\n'learning_rate\'] = [0.001, 0.01, 0.1, 1.0]\n# define the evaluation procedure\ncv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)\n# define the grid search procedure\ngrid_search = RandomizedSearchCV(estimator=AdaBoostClassifier(), param_distributions=grid, cv=cv, scoring=\n'accuracy\')\n# execute the grid search\ngrid_result = grid_search.fit(X_train, y_train)\n# summarize the best score and configuration\nprint("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))'
```

## 1.14 Testing Adaboost Classifier

```
[37]: adaboost = AdaBoostClassifier(n_estimators=162)
y_pred = train_model(X_train,X_test,y_train,adaboost)
evaluate("Adaboost",y_pred,y_test)
```


```
Adaboost accuracy: 0.8267696909272183
Adaboost f1 score: 0.8264669163545567
Adaboost precision: 0.8279139569784892
Adaboost recall: 0.825024925224327
```

```
/tmp/ipykernel_103467/1172794529.py:12: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
```

```
comparison_df = comparison_df.append(entry)
```

```
[38]: cf_matrix = confusion_matrix(y_test, y_pred)
print(cf_matrix)
ax= plt.subplot()
sns.heatmap(cf_matrix, annot=True,fmt='g',ax=ax,cmap='Blues')
ax.set_xlabel('Predicted labels');ax.set_ylabel('Actual labels');
ax.set_title('Confusion Matrix');
ax.xaxis.set_ticklabels(['Gamma', 'Hadron']); ax.yaxis.set_ticklabels(['Gamma', 'Hadron']);
```

```
[[1662  344]
 [ 351 1655]]
```



### 1.15 K-Nearest Neighbors (K-NN)

tunable parameters: `n_neighbors`


### 1.16 Hyperparameter tuning

```
[39]: k_neighbors = np.arange(5,100,2)
avg_accuracies = []
for k in k_neighbors:
    knn = KNeighborsClassifier(n_neighbors=k)
    t1 = time.time()
    acc_score = cross_validation(X_train,y_train,knn,cv=10)
    t2 = time.time() - t1
    print(f"k = {k}\naverage accuracy = {acc_score}\ntime = {t2}")
    avg_accuracies.append(acc_score)
```

```
k = 5    average accuracy = 0.8036102445520802    time = 0.9275436401367188
k = 7    average accuracy = 0.8073481925402952    time = 0.8415946960449219
k = 9    average accuracy = 0.8083120114203359    time = 0.8044657707214355
```

k = 11	average accuracy = 0.8079916126207483	time = 0.9410521984100342
k = 13	average accuracy = 0.8090591905426484	time = 0.9799323081970215
k = 15	average accuracy = 0.8117272801904605	time = 0.9461812973022461
k = 17	average accuracy = 0.811940271278585	time = 1.0473792552947998
k = 19	average accuracy = 0.8108725793357596	time = 1.0845611095428467
k = 21	average accuracy = 0.808844375119722	time = 1.0808305740356445
k = 23	average accuracy = 0.8082022092694452	time = 1.1431636810302734
k = 25	average accuracy = 0.8057466546260569	time = 1.0238127708435059
k = 27	average accuracy = 0.8065996451668811	time = 1.022817850112915
k = 29	average accuracy = 0.8062792463672933	time = 1.0328309535980225
k = 31	average accuracy = 0.8048913836667306	time = 1.0972936153411865
k = 33	average accuracy = 0.804784546059893	time = 1.015312910079956
k = 35	average accuracy = 0.8037174242216931	time = 1.0236971378326416
k = 37	average accuracy = 0.803824489870381	time = 1.06325101852417
k = 39	average accuracy = 0.8049983352944933	time = 1.052854061126709
k = 41	average accuracy = 0.803930757372593	time = 1.0777232646942139
k = 43	average accuracy = 0.8037177662844686	time = 1.0988304615020752
k = 45	average accuracy = 0.8037176522635434	time = 1.1346666812896729
k = 47	average accuracy = 0.8031824380410292	time = 1.0866870880126953
k = 49	average accuracy = 0.8016875096917786	time = 1.1202075481414795
k = 51	average accuracy = 0.8014741765408788	time = 1.1801459789276123
k = 53	average accuracy = 0.803076740643443	time = 1.1863205432891846
k = 55	average accuracy = 0.8032906438989682	time = 1.2490835189819336
k = 57	average accuracy = 0.8031840343339809	time = 1.1462140083312988
k = 59	average accuracy = 0.8037174242216933	time = 1.2071318626403809
k = 61	average accuracy = 0.8028631794506929	time = 1.2341325283050537
k = 63	average accuracy = 0.8024348028350163	time = 1.2632079124450684
k = 65	average accuracy = 0.8025410703372284	time = 1.2628123760223389
k = 67	average accuracy = 0.8035014685895154	time = 1.2377369403839111
k = 69	average accuracy = 0.8019007288217533	time = 1.2714812755584717
k = 71	average accuracy = 0.8008325807952277	time = 1.279567003250122
k = 73	average accuracy = 0.7993382225506025	time = 1.2782516479492188
k = 75	average accuracy = 0.7982706446287023	time = 1.32497239112854
k = 77	average accuracy = 0.7975225533389888	time = 1.3332364559173584
k = 79	average accuracy = 0.7967752601957512	time = 1.287491798400879
k = 81	average accuracy = 0.7964542912915379	time = 1.3423402309417725
k = 83	average accuracy = 0.7962414142243385	time = 1.4171881675720215
k = 85	average accuracy = 0.7958137217342127	time = 1.461984634399414
k = 87	average accuracy = 0.7954934369555502	time = 1.273763656616211
k = 89	average accuracy = 0.7952803318465005	time = 1.4719879627227783
k = 91	average accuracy = 0.7942119557781244	time = 1.3731598854064941
k = 93	average accuracy = 0.7950659725072746	time = 1.5490672588348389
k = 95	average accuracy = 0.7944254029499493	time = 1.4275591373443604
k = 97	average accuracy = 0.7945320125149368	time = 1.4569587707519531
k = 99	average accuracy = 0.7939987366481497	time = 1.4227516651153564

```
[40]: plt.plot(k_neighbors,np.array(avg_accuracies))
plt.show()
```




```
[41]: k_neighbors = np.arange(5,35,1)
avg_accuracies = []
for k in k_neighbors:
    knn = KNeighborsClassifier(n_neighbors=k)
    t1 = time.time()
    acc_score = cross_validation(X_train,y_train,knn,cv=10)
    t2 = time.time() - t1
    print(f"k = {k}\naverage accuracy = {acc_score}\ntime = {t2}")
    avg_accuracies.append(acc_score)
```

k = 5	average accuracy = 0.8036102445520802	time = 0.8367834091186523
k = 6	average accuracy = 0.794319705552363	time = 0.9571902751922607
k = 7	average accuracy = 0.8073481925402952	time = 0.9600963592529297
k = 8	average accuracy = 0.800088252196043	time = 0.8641960620880127
k = 9	average accuracy = 0.8083120114203359	time = 0.8481636047363281
k = 10	average accuracy = 0.8022234080398436	time = 0.9617252349853516
k = 11	average accuracy = 0.8079916126207483	time = 1.0233728885650635
k = 12	average accuracy = 0.8060689917813717	time = 0.9200780391693115



k = 13	average accuracy = 0.8090591905426484	time = 1.0408422946929932
k = 14	average accuracy = 0.8051068832152076	time = 1.063035011291504
k = 15	average accuracy = 0.8117272801904605	time = 1.0479013919830322
k = 16	average accuracy = 0.8089509846847092	time = 1.0557396411895752
k = 17	average accuracy = 0.811940271278585	time = 1.012247085571289
k = 18	average accuracy = 0.8074556002517582	time = 1.1496648788452148
k = 19	average accuracy = 0.8108725793357596	time = 1.009418249130249
k = 20	average accuracy = 0.8067088772131463	time = 1.030416488647461
k = 21	average accuracy = 0.808844375119722	time = 1.0960516929626465
k = 22	average accuracy = 0.8052132647383449	time = 0.958836555480957
k = 23	average accuracy = 0.8082022092694452	time = 0.9168310165405273
k = 24	average accuracy = 0.8044656295323318	time = 0.9686422348022461
k = 25	average accuracy = 0.8057466546260569	time = 0.974470853805542
k = 26	average accuracy = 0.8042516122558812	time = 1.1010453701019287
k = 27	average accuracy = 0.8065996451668811	time = 1.2187957763671875
k = 28	average accuracy = 0.8033970254221054	time = 1.363678216934204
k = 29	average accuracy = 0.8062792463672933	time = 1.1283354759216309
k = 30	average accuracy = 0.8037167400961425	time = 0.9869060516357422
k = 31	average accuracy = 0.8048913836667306	time = 1.0695619583129883
k = 32	average accuracy = 0.8019016409891544	time = 1.156188726425171
k = 33	average accuracy = 0.804784546059893	time = 1.0777318477630615
k = 34	average accuracy = 0.8019019830519298	time = 1.0036108493804932

```
[42]: plt.plot(k_neighbors,np.array(avg_accuracies))
      annot_max(k_neighbors,np.array(avg_accuracies))
      plt.show()
```



## 1.17 Testing KNN Classifier

```
[43]: knn = KNeighborsClassifier(n_neighbors=17)
      y_pred = train_model(X_train,X_test,y_train,knn)
      evaluate("KNN",y_pred,y_test)
```

```
KNN accuracy: 0.8103190428713859
KNN f1 score: 0.7971207677952545
KNN precision: 0.8567335243553008
KNN recall: 0.7452642073778664
```


```
/tmp/ipykernel_103467/1172794529.py:12: FutureWarning: The frame.append method
is deprecated and will be removed from pandas in a future version. Use
pandas.concat instead.
```

```
    comparison_df = comparison_df.append(entry)
```

```
[44]: cf_matrix = confusion_matrix(y_test, y_pred)
      print(cf_matrix)
      ax= plt.subplot()
      sns.heatmap(cf_matrix, annot=True,fmt='g',ax=ax,cmap='Blues')
      ax.set_xlabel('Predicted labels');ax.set_ylabel('Actual labels');
      ax.set_title('Confusion Matrix');
```

```
ax.xaxis.set_ticklabels(['Gamma', 'Hadron']); ax.yaxis.set_ticklabels(['Gamma', 'Hadron']);
```

```
[[1756  250]
 [ 511 1495]]
```



## 1.18 Random Forest

A random forest is a meta estimator that fits a number of decision tree classifiers on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control overfitting. The sub-sample size is controlled with the `max_samples` parameter if `bootstrap=True` (default), otherwise the whole dataset is used to build each tree.

tunable parameters: `n_estimators`


`n_estimators`: The number of trees in the forest.

## 1.19 Hyperparameter tuning

```
[45]: n_estimators = np.linspace(50,300,20,dtype=int)
avg_accuracies = []
for n in n_estimators:
    rf = RandomForestClassifier(n_estimators=n)
    t1 = time.time()
    acc_score = cross_validation(X_train,y_train,rf,cv=10)
    t2 = time.time() - t1
    print(f"n = {n}\naverage accuracy = {acc_score}\ntime = {t2}")
    avg_accuracies.append(acc_score)
```

n = 50	average accuracy = 0.8546554743726569	time = 14.546619415283203
n = 63	average accuracy = 0.8533742212370814	time = 18.63932204246521
n = 76	average accuracy = 0.8547642503352215	time = 22.82245397567749
n = 89	average accuracy = 0.8529476689562067	time = 26.597578048706055
n = 102	average accuracy = 0.852735590035483	time = 29.100209712982178
n = 115	average accuracy = 0.8547636802305959	time = 30.842748403549194
n = 128	average accuracy = 0.8567916564047833	time = 34.293354511260986
n = 142	average accuracy = 0.8550833949046328	time = 38.108360052108765
n = 155	average accuracy = 0.8578606025777852	time = 41.908435583114624
n = 168	average accuracy = 0.8538020277481323	time = 45.20749306678772
n = 181	average accuracy = 0.8546560444772824	time = 48.60418939590454
n = 194	average accuracy = 0.8570068138904849	time = 52.40931987762451
n = 207	average accuracy = 0.8558319422780469	time = 59.73830986022949
n = 221	average accuracy = 0.857005673681234	time = 59.9186155796051
n = 234	average accuracy = 0.85636624433316	time = 60.4331157207489
n = 247	average accuracy = 0.8572191208530588	time = 63.94174265861511
n = 260	average accuracy = 0.8561514289102335	time = 67.44234561920166
n = 273	average accuracy = 0.8555110873947587	time = 70.46408534049988
n = 286	average accuracy = 0.855297982285709	time = 74.47242665290833
n = 300	average accuracy = 0.8547637942515209	time = 77.95634722709656


```
[46]: plt.plot(n_estimators,np.array(avg_accuracies))
plt.show()
```



```
[47]: n_estimators = np.linspace(125,160,10,dtype=int)
avg_accuracies = []
for n in n_estimators:
    rf = RandomForestClassifier(n_estimators=n)
    t1 = time.time()
    acc_score = cross_validation(X_train,y_train,rf,cv=10)
    t2 = time.time() - t1
    print(f"n = {n}\taverage accuracy = {acc_score}\tttime = {t2}")
    avg_accuracies.append(acc_score)
```

n = 125 average accuracy = 0.8547644783770716	time = 32.737011432647705
n = 128 average accuracy = 0.8561508588056078	time = 32.50386452674866
n = 132 average accuracy = 0.8548707458792839	time = 34.80695652961731
n = 136 average accuracy = 0.8522000337501938	time = 35.76816487312317
n = 140 average accuracy = 0.8544441935984091	time = 38.21756052970886
n = 144 average accuracy = 0.8563655602076092	time = 42.68096613883972
n = 148 average accuracy = 0.8538029399155332	time = 44.07245111465454
n = 152 average accuracy = 0.8542298342591833	time = 41.50951528549194
n = 156 average accuracy = 0.854657640770234	time = 42.76413345336914
n = 160 average accuracy = 0.8559387798848844	time = 45.390727519989014

```
[48]: plt.plot(n_estimators,np.array(avg accuracies))
plt.show()
```




```
[49]: n_estimators = np.linspace(140,145,5,dtype=int)
avg accuracies = []
for n in n_estimators:
    rf = RandomForestClassifier(n_estimators=n)
    t1 = time.time()
    acc_score = cross_validation(X_train,y_train,rf,cv=10)
    t2 = time.time() - t1
    print(f"n = {n}\taverage accuracy = {acc_score}\ttime = {t2}")
    avg accuracies.append(acc_score)
```

```
n = 140 average accuracy = 0.8531625984000583    time = 37.9939911365509
n = 141 average accuracy = 0.8552982103275593    time = 40.81869029998779
n = 142 average accuracy = 0.8555108593529083    time = 39.429386138916016
n = 143 average accuracy = 0.8559375256547082    time = 38.95640182495117
n = 145 average accuracy = 0.854655702414507     time = 40.980098724365234
```

```
[50]: plt.plot(n_estimators,np.array(avg accuracies))
annot_max(n_estimators,np.array(avg accuracies))
```

```
plt.show()
```



best n\_estimators value: ~143

## 1.20 Testing Random Forest Classifier

```
[51]: rf = RandomForestClassifier(n_estimators=143)
y_pred = train_model(X_train,X_test,y_train,rf)
evaluate("Random Forest",y_pred,y_test)
```

```
Random Forest accuracy: 0.8579262213359921
Random Forest f1 score: 0.854368932038835
Random Forest precision: 0.8763102725366876
Random Forest recall: 0.8334995014955134
```


```
/tmp/ipykernel_103467/1172794529.py:12: FutureWarning: The frame.append method
is deprecated and will be removed from pandas in a future version. Use
pandas.concat instead.
```

```
comparison_df = comparison_df.append(entry)
```

```
[52]: cf_matrix = confusion_matrix(y_test, y_pred)
print(cf_matrix)
ax= plt.subplot()
```

```
sns.heatmap(cf_matrix, annot=True,fmt='g',ax=ax,cmap='Blues')
ax.set_xlabel('Predicted labels');ax.set_ylabel('Actual labels');
ax.set_title('Confusion Matrix');
ax.xaxis.set_ticklabels(['Gamma', 'Hadron']); ax.yaxis.set_ticklabels(['Gamma', 'Hadron']);
```

```
[[1770  236]
 [ 334 1672]]
```



## 1.21 Naïve Bayes

Tunable parameters: None

## 1.22 Choosing var\_smoothing parameter

var\_smoothing: Portion of the largest variance of all features that is added to variances for calculation stability.

```
[53]: var_smoothing = np.logspace(0,-9, num=10)
      avg_accuracies = []
      for v in var_smoothing:
```



```

nb = GaussianNB(var_smoothing=v)
t1 = time.time()
acc_score = cross_validation(X_train,y_train,nb,cv=10)
t2 = time.time() - t1
print(f"v = {v}\naverage accuracy = {acc_score}\ttime = {t2}")
avg_accuracies.append(acc_score)

```

```

v = 1.0 average accuracy = 0.6869899844019376    time = 0.04506111145019531
v = 0.1 average accuracy = 0.6889132893668646    time = 0.043242692947387695
v = 0.01          average accuracy = 0.6598669147761999    time =
0.04018568992614746
v = 0.001          average accuracy = 0.651217173375658    time =
0.03883981704711914
v = 0.0001          average accuracy = 0.650790279032008    time =
0.03502988815307617
v = 1e-05          average accuracy = 0.6506836694670207    time =
0.04058432579040527
v = 1e-06          average accuracy = 0.6506836694670207    time =
0.05889701843261719
v = 1e-07          average accuracy = 0.6506836694670207    time =
0.05028414726257324
v = 1e-08          average accuracy = 0.6506836694670207    time =
0.046877384185791016
v = 1e-09          average accuracy = 0.6506836694670207    time =
0.03712105751037598

```

Best var\_smoothing value: ~0.1

## 1.23 Feature Selection

```

[54]: from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
def feature_select(X_train,y_train):
    n_samples = X_train.shape[0]
    for i in range(1,X.shape[1]+1):
        select = SelectKBest(k=i, score_func=chi2)
        select.fit(X_train, y_train)
        X_new = select.transform(X_train)
        acc_score = cross_validation(X_new,y_train,GaussianNB(var_smoothing=0.
↪1),cv=10)
        print(f"no. of features= {i}\naverage accuracy = {acc_score}")
        avg_accuracies.append(acc_score)

```

```

[55]: feature_select(X_train,y_train)

```

no. of features= 1	average accuracy = 0.7047201242371999
no. of features= 2	average accuracy = 0.7367549872752648
no. of features= 3	average accuracy = 0.723832083663994
no. of features= 4	average accuracy = 0.7236181804084686
no. of features= 5	average accuracy = 0.7129420591266908
no. of features= 6	average accuracy = 0.7108070173038156
no. of features= 7	average accuracy = 0.70140895657171
no. of features= 8	average accuracy = 0.6986320909613333
no. of features= 9	average accuracy = 0.6968163077287943
no. of features= 10	average accuracy = 0.6889132893668646

The best 2 features yield max accuracy

## 1.24 Testing Naive Bayes Classifier

```
[56]: select = SelectKBest(k=2, score_func=chi2)
select.fit(X_train, y_train)
X_new = select.transform(X_train)
X_test_new = select.transform(X_test)
y_pred = train_model(X_new,X_test_new,y_train,GaussianNB(var_smoothing=0.1))
evaluate("Naive Bayes",y_pred,y_test)
```

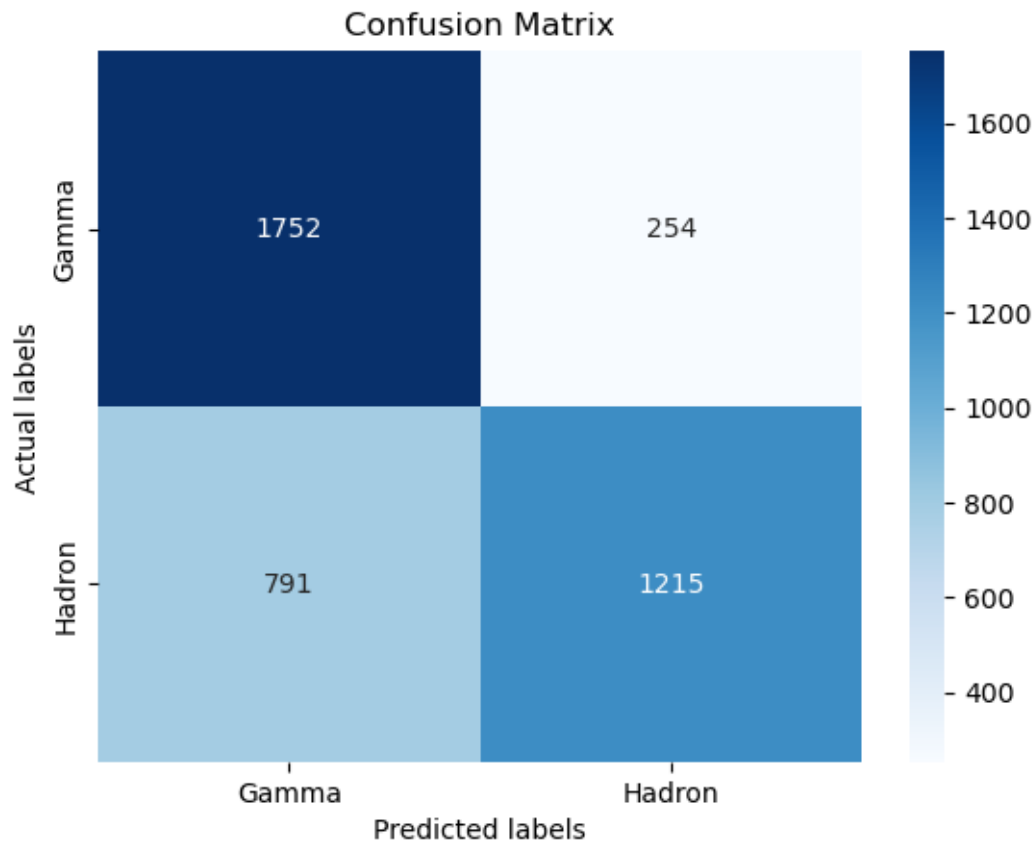
Naive Bayes accuracy: 0.7395314057826521  
Naive Bayes f1 score: 0.6992805755395683  
Naive Bayes precision: 0.8270932607215793  
Naive Bayes recall: 0.6056829511465603

/tmp/ipykernel\_103467/1172794529.py:12: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
comparison_df = comparison_df.append(entry)
```


```
[57]: cf_matrix = confusion_matrix(y_test, y_pred)
print(cf_matrix)
ax= plt.subplot()
sns.heatmap(cf_matrix, annot=True,fmt='g',ax=ax,cmap='Blues')
ax.set_xlabel('Predicted labels');ax.set_ylabel('Actual labels');
ax.set_title('Confusion Matrix');
ax.xaxis.set_ticklabels(['Gamma', 'Hadron']); ax.yaxis.set_ticklabels(['Gamma', 'Hadron']);
```

```
[[1752  254]
 [ 791 1215]]
```



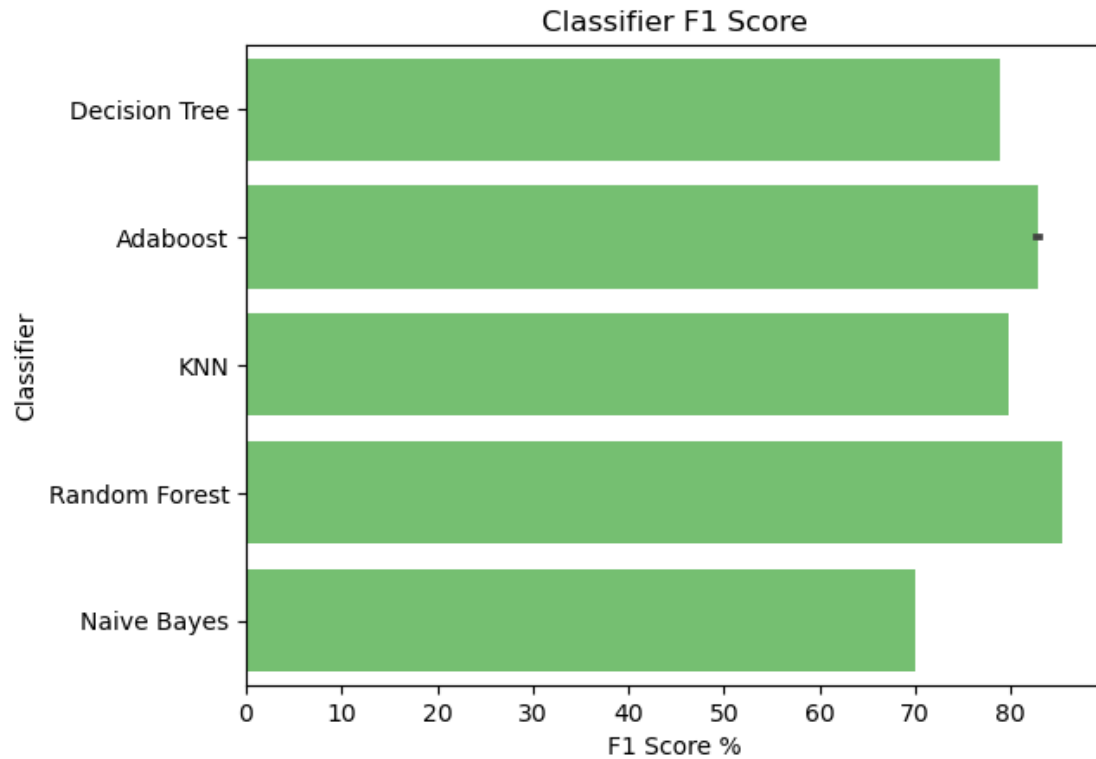
```
[58]: sns.set_color_codes("muted")
sns.barplot(x='Accuracy', y='Classifier', data=comparison_df, color="b")

plt.xlabel('Accuracy %')
plt.title('Classifier Accuracy')
plt.show()
```



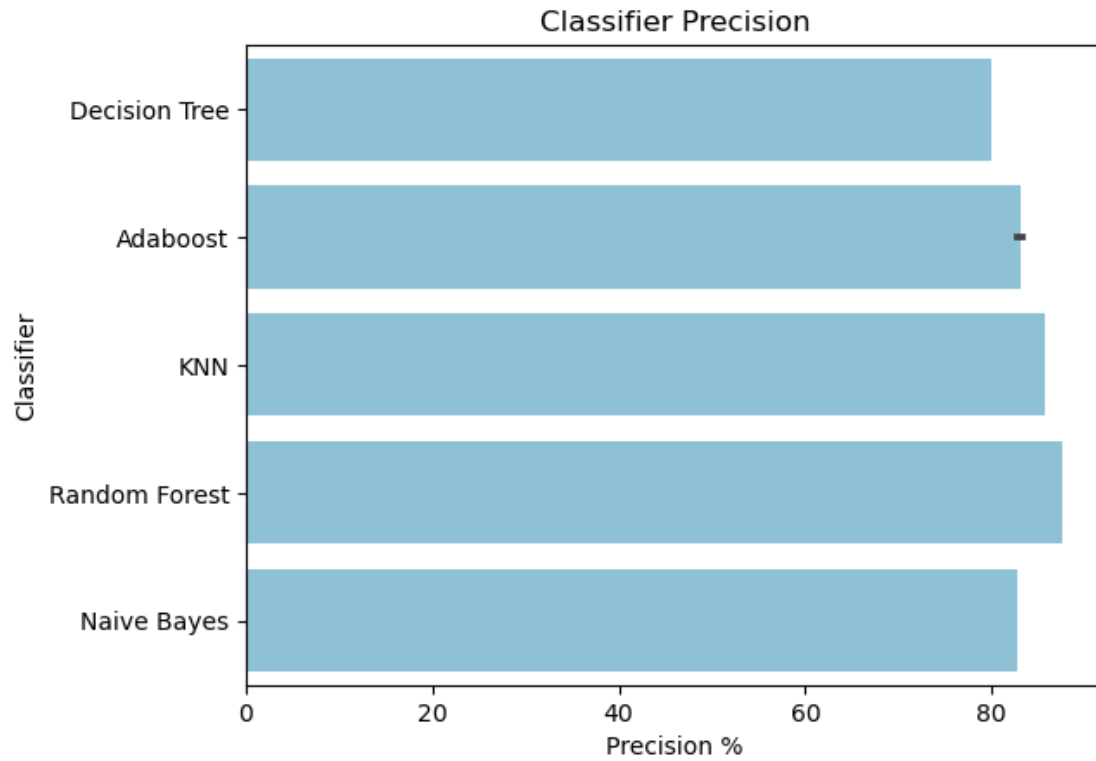
```
[59]: sns.set_color_codes("muted")
sns.barplot(x='F1 score', y='Classifier', data=comparison_df, color="g")

plt.xlabel('F1 Score %')
plt.title('Classifier F1 Score')
plt.show()
```



```
[60]: sns.set_color_codes("muted")
sns.barplot(x='Precision', y='Classifier', data=comparison_df, color="c")

plt.xlabel('Precision %')
plt.title('Classifier Precision')
plt.show()
```



```
[61]: sns.set_color_codes("muted")
sns.barplot(x='Recall', y='Classifier', data=comparison_df, color="r")

plt.xlabel('Recall %')
plt.title('Classifier Recall')
plt.show()
```



[ ]: